

Unit 1: Introduction to R Programming

Q1. What is R, and what is its primary use in data analysis?

R is a powerful open-source programming language and software environment primarily used for statistical computing, data analysis, and graphical representation. Developed by Ross Ihaka and Robert Gentleman in 1993 at the University of Auckland, New Zealand, R is a free implementation of the S programming language and incorporates lexical scoping inspired by Scheme.

The primary use of R is in data analysis and statistical computing. It provides a vast collection of built-in functions and packages that enable users to perform tasks such as descriptive statistics, regression analysis, hypothesis testing, classification, clustering, time-series forecasting, and machine learning. One of R's major strengths is its ability to produce high-quality, customizable graphical representations of data using libraries like **ggplot2**, **lattice**, and **plotly**.

R supports various data types and structures such as vectors, matrices, data frames, and lists. It is widely used in academia, research, business analytics, and industries dealing with large datasets due to its flexibility and performance.

Furthermore, R can integrate with other programming languages such as C, C++, Java, and Python, making it highly versatile. As it is platform-independent and open-source, it is freely available for all operating systems without licensing costs. The strong and active user community contributes to constant enhancements and support.

In conclusion, R is a comprehensive tool for statistical computing and data analysis, offering vast capabilities for data manipulation, visualization, and modeling, making it an essential language in the field of data science.

Q2. Describe the key features that make R a preferred tool for statistical computing.

R has become one of the most preferred tools for statisticians and data scientists due to its rich set of features tailored for statistical analysis and visualization. The key features include:

1. **Comprehensive Statistical Techniques:**

R supports a wide range of statistical methods such as linear and nonlinear modeling, hypothesis testing, time-series analysis, classification, clustering, and more.

2. **Advanced Data Visualization:**

R provides powerful graphical capabilities. Libraries like **ggplot2**, **lattice**, and **shiny** allow users to create visually appealing and interactive charts and dashboards.

3. **Extensive Package Ecosystem:**

The CRAN repository hosts over 15,000 packages that extend R's functionality in areas such as bioinformatics, econometrics, text mining, and geospatial analysis.

4. Open Source and Free:

R is freely available for use and distribution. Its open-source nature encourages collaboration and constant improvements from a global community of developers.

5. Platform Independent:

R can be used on Windows, macOS, and Linux systems, providing flexibility across different computing environments.

6. Integration with Other Languages:

R can easily integrate with C, C++, Python, Java, and SQL, allowing seamless data exchange and performance optimization.

7. Efficient Data Handling:

R efficiently handles various data structures, including large datasets, using data frames, lists, and matrices.

8. Active Community and Documentation:

R has a strong global user base that contributes to forums, tutorials, packages, and comprehensive documentation.

9. Reproducible Research Support:

With tools like R Markdown and Knitr, R enables the creation of dynamic reports that include code, results, and visualizations.

Q3. What is RStudio? List its main components.

RStudio is an Integrated Development Environment (IDE) specifically designed for the R programming language. It provides a user-friendly interface and a set of tools that help users write, debug, and manage R code efficiently. RStudio is available in both desktop and server versions, making it accessible for both personal and enterprise-level applications.

The primary function of RStudio is to make working with R more efficient by providing a structured environment for writing scripts, visualizing data, managing projects, and exploring R objects. Once R is installed on the system, RStudio serves as a powerful interface to interact with it more intuitively.

Main Components of RStudio:

1. Console Panel (Left Bottom):

- This is where users type and execute R commands directly.
- The output is displayed immediately.

2. Source/Script Editor (Top Left):

- Used to write and edit R scripts (.R files), R Markdown documents (.Rmd), or Shiny applications.
- Supports syntax highlighting and code completion.

3. Environment/History Panel (Top Right):

- **Environment Tab:** Shows all active objects, variables, and datasets in the current R session.
- **History Tab:** Displays a log of all commands executed in the console.

4. Files/Plots/Packages/Help/Viewer Panel (Bottom Right):

- **Files Tab:** Shows the files and folders in the current working directory.
- **Plots Tab:** Displays plots and charts generated from R commands.
- **Packages Tab:** Allows users to manage installed packages and install new ones.
- **Help Tab:** Offers documentation and help related to functions and packages.
- **Viewer Tab:** Displays web content such as HTML outputs from R Markdown.

These components make RStudio a comprehensive and efficient environment for programming in R.

Q4. Describe the process of quitting RStudio.

Quitting RStudio involves closing the current R session, which also prompts the user to save the current workspace. This is important because if the workspace is saved, all variables and loaded objects will be reloaded the next time RStudio is opened.

Steps to Quit RStudio:

1. Manual Exit via GUI:

- Click the **File** menu from the top-left corner.
- Select **Quit Session** or click the close (X) button on the top-right corner of the RStudio window.
- A prompt will appear asking whether to save the current workspace image (`.RData`).
- Choose "Save", "Don't Save", or "Cancel" as per requirement.

2. Using Console Command:

- You can also use the function `q()` or `quit()` in the console.
- Syntax: `q()`
- This will prompt the same save workspace dialog.

3. Auto-Save Option:

- Users can also configure RStudio settings to always save or never save the workspace when quitting.

- Navigate to: **Tools > Global Options > General > Save workspace to .RData on exit** and set preferences.

Note: It is recommended to save scripts using .R files to ensure reproducibility, rather than relying only on the workspace image.

Q5. Explain the difference between R and RStudio. How do they interact with each other?

R and RStudio are two essential tools used in the field of statistical computing and data analysis, but they serve different purposes. Understanding their differences and how they interact is fundamental for effective usage.

R (The Language and Engine): R is a programming language and software environment used for statistical analysis, data visualization, and data manipulation. It provides the core functionality and computational capabilities. R runs in the background and processes all commands and computations.

RStudio (The Interface): RStudio, on the other hand, is an Integrated Development Environment (IDE) for R. It provides a user-friendly interface that makes it easier to write R code, manage files and projects, and visualize outputs. It includes features such as syntax highlighting, debugging tools, project management, and integrated help documentation.

Interaction Between R and RStudio:

- RStudio acts as a front-end to the R engine.
- When users type code into RStudio's console or script editor and run it, RStudio sends this code to the R engine for execution.
- The results computed by R are then displayed in RStudio's console or plot panel.

Analogy: Think of R as the engine of a car and RStudio as the dashboard. The engine (R) does all the heavy lifting, while the dashboard (RStudio) makes it easy for the driver (user) to control and monitor the vehicle.

Conclusion: While R can be used on its own, RStudio significantly enhances productivity, especially for beginners and professionals working on large projects. Together, they offer a complete environment for statistical analysis and data science workflows.

Unit 2: R Data Structures and Manipulation

Q1. How do you create a variable in R? Provide an example.

In R, a variable is used to store data such as numbers, text, or logical values. You can create a variable using the assignment operators `<-`, `=`, or `->`, though `<-` is most commonly used in R programming.

Variables do not require explicit declaration of data type, as R is a dynamically typed language.

Syntax and Example:

```
x <- 10      # Assigning numeric value to x
y = "Hello"  # Assigning character string to y
20 -> z      # Assigning value to z using right assignment
```

Here, x stores a numeric value 10, y stores a character string, and z stores the number 20. These variables can be used in computations, printed, or manipulated further. R automatically determines the data type of the variable based on the value assigned.

Q2. Explain the use of conditional statements in R with an example.

Conditional statements in R control the flow of execution depending on certain conditions. The most common conditional statements in R are `if`, `if...else`, and `if...else if...else`. They allow the execution of different blocks of code based on whether a condition is TRUE or FALSE.

Example 1: Simple if statement

```
x <- 5
if (x > 0) {
  print("x is positive")
}
```

Example 2: if...else statement

```
x <- -2
if (x > 0) {
  print("Positive")
} else {
  print("Negative or Zero")
}
```

Example 3: if...else if...else

```
x <- 0
if (x > 0) {
  print("Positive")
} else if (x < 0) {
  print("Negative")
} else {
  print("Zero")
}
```

These structures make decision-making in programs flexible and dynamic, allowing responses to varying inputs.

Q3. Write an R function to calculate the sum of elements in a vector.

Functions in R allow users to encapsulate logic and reuse code efficiently. Below is an R function that calculates the sum of all elements in a numeric vector:

Example:

```
# Define the function
sum_vector <- function(vec) {
  total <- sum(vec)
  return(total)
}
```

```
# Call the function
my_vector <- c(2, 4, 6, 8)
result <- sum_vector(my_vector)
print(result) # Output: 20
```

Explanation:

- `sum()` is a built-in R function that calculates the total of numeric values.
- The user-defined function `sum_vector()` takes a vector `vec` as input and returns the sum of its elements.

This demonstrates how R supports modular and functional programming.

Unit 3: R Packages and Functions**Q1. What is an R package? Why is it used?**

An R package is a collection of R functions, data, and documentation that are grouped together for reuse and sharing. Packages are essential in R as they extend the functionality of the base R installation. They allow users to perform specific tasks such as data visualization, statistical modeling, machine learning, and more.

Purpose and Use:

- Packages make it easier to organize and reuse code.
- They offer pre-built functions to perform complex tasks easily.
- CRAN (Comprehensive R Archive Network) hosts thousands of R packages.

Example packages include `ggplot2` for plotting, `dplyr` for data manipulation, and `caret` for machine learning.

Q2. How do you install a package in RStudio?

To install a package in RStudio, you can use the `install.packages()` function in the console. This downloads the package from CRAN and installs it locally on your machine.

Syntax:

```
install.packages("ggplot2") # Installing ggplot2 package
```

After installation, the package must be loaded into the session using:

```
library(ggplot2)
```

Alternatively, you can install packages through RStudio's graphical interface:

- Go to **Tools > Install Packages**, enter the package name, and click install.
-

Q3. Write a simple R function to calculate the square of a number.

Example:

```
square_number <- function(x) {  
  return(x^2)  
}
```

```
# Usage  
square_number(4) # Output: 16
```

Explanation:

- The function `square_number()` takes an input `x` and returns its square using the exponent operator `^`.
 - This is an example of a user-defined function for modular and reusable programming in R.
-

Q4. Explain how to download and import data in R.

To work with external data in R, you can use functions such as `read.csv()` for CSV files or `read.table()` for text files. You can import data either from local storage or directly from a URL.

Example: Import from URL

```
data <- read.csv("https://people.sc.fsu.edu/~jburkardt/data/csv/airtravel.csv")  
head(data)
```

Example: Import from local file

```
data <- read.csv("C:/Users/Username/Documents/myfile.csv")
```

Functions for Importing Data:

- `read.csv()`: Reads CSV files.
- `read.table()`: Reads tabular text files.
- `read_excel()`: Reads Excel files (requires `readxl` package).

These functions allow R to read structured datasets into data frames for further analysis.

Q5. Describe the process of creating a custom function in R. Provide an example.

Creating custom functions in R allows users to wrap a set of instructions in a single reusable block. Functions are created using the `function()` keyword.

Syntax:

```
function_name <- function(parameters) {  
  # code block  
  return(result)  
}
```

Example:

```
multiply_numbers <- function(a, b) {  
  product <- a * b  
  return(product)  
}  
  
# Calling the function  
result <- multiply_numbers(5, 4)  
print(result) # Output: 20
```

Explanation:

- The function `multiply_numbers()` takes two arguments `a` and `b`.
- It returns their product using the `*` operator.
- The `return()` function sends back the result to the calling environment.

Functions in R enhance code readability, reduce redundancy, and allow structured programming.

Unit 4: Matrices, Arrays, and Lists**Q1. How do you create a matrix in R? Provide an example.**

A matrix in R is a two-dimensional data structure that contains elements of the same data type (either numeric, character, or logical). You can create a matrix using the `matrix()` function.

Syntax and Example:

```
mat <- matrix(1:6, nrow = 2, ncol = 3)  
print(mat)
```

Output:

```
      [,1] [,2] [,3]  
[1,]    1    3    5  
[2,]    2    4    6
```

Explanation:

- The numbers 1 to 6 are filled into a 2-row and 3-column matrix by column-wise filling (default in R).

Q2. Describe the operations that can be performed on matrices in R.

Matrices in R support various mathematical and indexing operations. These include:

1. Arithmetic Operations:

- Addition: `mat1 + mat2`
- Subtraction: `mat1 - mat2`
- Multiplication (element-wise): `mat1 * mat2`

- Matrix multiplication: `mat1 %*% mat2`

2. Transposition:

- `t(mat)` transposes the matrix.

3. Accessing Elements:

- `mat[1, 2]` accesses element at row 1, column 2.
- `mat[, 2]` accesses the entire second column.

4. Functions on Matrices:

- `rowSums(mat)`, `colSums(mat)`
- `rowMeans(mat)`, `colMeans(mat)`

5. Binding Matrices:

- `rbind(mat1, mat2)` – Row bind.
- `cbind(mat1, mat2)` – Column bind.

These operations enable advanced data analysis and linear algebra tasks in R.

Q3. What is the difference between a vector and a matrix in R?

Feature	Vector	Matrix
Dimensions	1D	2D
Homogeneity	Must be same data type	Must be same data type
Access Method	Single index	Two indices [row, column]
Creation	<code>c()</code>	<code>matrix()</code>
Example	<code>v <- c(1, 2, 3)</code>	<code>m <- matrix(1:6, 2, 3)</code>

Explanation: A vector is a one-dimensional sequence of elements, whereas a matrix is a two-dimensional structure organized in rows and columns. Both are homogeneous (same data type), but matrices support more complex operations such as matrix multiplication.

Q4. How do you add or delete rows and columns in a matrix?

Adding Rows and Columns:

- To add a row: Use `rbind()`

```
m <- matrix(1:6, nrow=2)
m_new <- rbind(m, c(7, 8, 9))
```

- To add a column: Use `cbind()`

```
m_new <- cbind(m, c(10, 11))
```

Deleting Rows and Columns:

- To delete a row:

```
m <- m[-2, ] # deletes 2nd row
```

- To delete a column:

```
m <- m[, -1] # deletes 1st column
```

Explanation: The minus (-) sign in indexing is used to exclude rows or columns. `rbind()` and `cbind()` are used to add rows or columns by combining vectors or matrices.

Q5. How would you handle higher-dimensional arrays in R? Explain with an example.

An array in R can have more than two dimensions. Arrays are created using the `array()` function and are useful for storing data in three or more dimensions.

Syntax and Example:

```
arr <- array(1:24, dim = c(3, 4, 2))  
print(arr)
```

Explanation:

- This creates a 3x4x2 array, meaning two 3x4 matrices stacked together.
- Elements are filled by columns across the array.

Accessing Elements:

```
arr[1, 2, 1] # Accesses element at row 1, column 2, first matrix
```

Applications: Higher-dimensional arrays are useful in image processing, time-series data, simulations, and scientific computing where data has more than two axes (e.g., x, y, time).

Unit 5: Data Frames

Q1. What is a data frame in R? How is it different from a matrix?

A data frame is a two-dimensional, tabular data structure in R where each column can contain different types of data (numeric, character, logical, etc.). It is similar to a spreadsheet or SQL table.

Difference from Matrix:

- A matrix can only hold one data type (e.g., all numeric or all character), whereas a data frame can have mixed types.
- Data frames are more suitable for real-world datasets with varied data types.

Example:

```
df <- data.frame(Name = c("John", "Alice"), Age = c(23, 25), Passed = c(TRUE,  
FALSE))
```

```
print(df)
```

Q2. How do you create a data frame in R? Provide an example.

You can create a data frame using the `data.frame()` function. You specify column names and vectors of equal length.

Example:

```
names <- c("Sam", "Lily")
ages <- c(22, 24)
grades <- c("A", "B")
df <- data.frame(Name = names, Age = ages, Grade = grades)
print(df)
```

Explanation: This creates a data frame with three columns: Name, Age, and Grade, and two rows of student data.

Q3. What are factors in R? Explain their importance.

Factors in R are used to represent categorical data, such as gender, status, or level. Internally, factors are stored as integers with corresponding character labels.

Importance:

- Factors help in efficient storage and analysis of categorical data.
- They are essential for statistical modeling, such as regression or classification.

Example:

```
gender <- factor(c("Male", "Female", "Male", "Female"))
print(gender)
```

Output:

```
[1] Male    Female Male    Female
Levels: Female Male
```

Q4. How can you merge two data frames in R?

R allows merging of data frames using the `merge()` function. You can merge by common columns or row names.

Example:

```
df1 <- data.frame(ID = c(1, 2), Name = c("Alex", "Brian"))
df2 <- data.frame(ID = c(1, 2), Score = c(85, 90))
merged_df <- merge(df1, df2, by = "ID")
print(merged_df)
```

Output:

	ID	Name	Score
1	1	Alex	85
2	2	Brian	90

Explanation: `merge()` joins the two data frames on the common column `ID`, combining matching rows.

Q5. What is the `apply()` function in R, and how does it work with data frames?

The `apply()` function is used to apply a function to the rows or columns of a matrix or data frame.

Syntax:

```
apply(X, MARGIN, FUN)
```

- `X`: Data frame or matrix
- `MARGIN`: 1 for rows, 2 for columns
- `FUN`: Function to apply

Example:

```
df <- data.frame(A = c(1, 2, 3), B = c(4, 5, 6))  
apply(df, 2, mean) # Apply mean to columns
```

Output:

A	B
2	5

Explanation: Here, `apply()` calculates the mean of each column. It is commonly used for summarization and transformations.

Q6. Write an R function to calculate the mean of each column in a data frame.**Example:**

```
mean_columns <- function(df) {  
  return(apply(df, 2, mean))  
}  
  
# Usage  
mydata <- data.frame(x = c(1, 2, 3), y = c(4, 5, 6))  
result <- mean_columns(mydata)  
print(result)
```

Explanation: The function `mean_columns()` applies the `mean` function to each column using `apply()` with `MARGIN = 2`.

Q7. How can you deal with scope issues when working with functions and objects in R?

Scope issues in R arise when variables with the same name exist in multiple environments (e.g., inside and outside a function).

Solutions:

1. Use of Local and Global Environments:

- Variables defined inside functions are local.
- Use `<<-` to assign to a variable in the global scope.

2. Avoid Name Conflicts:

- Use unique and descriptive variable names.

3. Accessing Global Variables:

- Global variables can be accessed inside a function without redefining unless shadowed by local variables.

Example:

```
a <- 10
myfunc <- function() {
  a <- 5
  print(a) # prints local a = 5
}
myfunc()
print(a) # prints global a = 10
```

Q8. Write a script to work with tables in R, including creating and applying functions.

Example Script:

```
# Create a data frame
data <- data.frame(Name = c("Tom", "Jerry", "Spike"), Marks = c(80, 90, 85))

# Create a table of frequencies (Marks > 85)
table_result <- table(data$Marks > 85)
print(table_result)

# Define a function to classify grades
grade_function <- function(score) {
  if (score >= 90) return("A")
  else if (score >= 80) return("B")
  else return("C")
}

# Apply function to create a new column
data$Grade <- apply(data$Marks, grade_function)
print(data)
```

Explanation:

- A table is created using `table()` to check how many scores are above 85.
- A custom function `grade_function()` is defined and applied to assign grades based on marks.